

# Keras Notes

```
[2]: import tensorflow as tf

[3]: print(tf.__version__)
2.16.2

[5]: # create a constant tensor A
A= tf.constant([[4,2],[6,1]])
A

[5]: <tf.Tensor: shape=(2, 2), dtype=int32, numpy=
array([[4, 2],
       [6, 1]])>

[9]: V=tf.Variable([[9,2],[6,5]])
V

[9]: <tf.Variable 'Variable:0' shape=(2, 2) dtype=int32, numpy=
array([[9, 2],
       [6, 5]])>

[17]: AV=tf.concat(values=[A,V],axis=0)
av=tf.concat(values=[A,V],axis=1)
print(AV, '\n\n',av)

tf.Tensor(
[[9 2]
 [6 5]
 [9 2]
 [6 5]], shape=(4, 2), dtype=int32)

tf.Tensor(
[[9 2 9 2]
 [6 5 6 5]], shape=(2, 4), dtype=int32)

[19]: zero_tflw=tf.zeros(shape=[3,4],dtype=tf.int32)
zero_tflw

[19]: <tf.Tensor: shape=(3, 4), dtype=int32, numpy=
array([[0, 0, 0, 0],
       [0, 0, 0, 0],
       [0, 0, 0, 0]])>
```

```
[28]: one_tflw=tf.ones(shape=[3,4],dtype=tf.float32)
one_tflw
```

```
[28]: <tf.Tensor: shape=(3, 4), dtype=float32, numpy=
array([[1., 1., 1., 1.],
       [1., 1., 1., 1.],
       [1., 1., 1., 1.]], dtype=float32)>
```

```
[68]: # identity matrix
IM=tf.constant([[1,2],[3,4],[5,6]])
rows, columns= IM.shape
print('Row:',rows,'\\nColumns:',columns)

IM_indent=tf.eye(num_rows=rows,num_columns=columns, dtype=tf.int32)
print('indentity matrix of IM is:\\n',IM_indent)
```

```
Row: 3
Columns: 2
indentity matrix of IM is:
tf.Tensor(
[[1 0]
 [0 1]
 [0 0]], shape=(3, 2), dtype=int32)
```

```
[34]: random_generate=tf.random.uniform(shape=[2,4],dtype=tf.float32)
random_generate
```

```
[34]: <tf.Tensor: shape=(2, 4), dtype=float32, numpy=
array([0.7430968 , 0.52440953, 0.8917074 , 0.6367301 ],
      [0.8461467 , 0.8472326 , 0.55651677, 0.8393111 ]], dtype=float32)>
```

```
[35]: reshape_tensor=tf.reshape(tensor= random_generate
                               ,shape= [4,2])
reshape_tensor
```

```
[35]: <tf.Tensor: shape=(4, 2), dtype=float32, numpy=
array([[0.7430968 , 0.52440953],
       [0.8917074 , 0.6367301 ],
       [0.8461467 , 0.8472326 ],
       [0.55651677, 0.8393111 ]], dtype=float32)>
```

```
[39]: # typecast a tensor
x=tf.constant([[9,2],[6,5]],dtype=tf.float32)
print(x)

X=tf.cast(x,tf.int32)
print('\n',X)

tf.Tensor(
[[9. 2.]
 [6. 5.]], shape=(2, 2), dtype=float32)

tf.Tensor(
[[9 2]
 [6 5]], shape=(2, 2), dtype=int32)

[43]: t=tf.transpose(X)
print('transpose of X:\n',t)

transpose of X:
tf.Tensor(
[[9 6]
 [2 5]], shape=(2, 2), dtype=int32)

[49]: A=tf.constant([[2,3],[4,5]])
V=tf.constant([[4],[2]])

[52]: # matmul matrix multiplier
AV=tf.matmul(A,V)
AV

[52]: <tf.Tensor: shape=(2, 1), dtype=int32, numpy=
array([[14],
       [26]])>

[54]: # Element wise multiply
av=tf.multiply(A,V)
av

[54]: <tf.Tensor: shape=(2, 2), dtype=int32, numpy=
array([[ 8, 12],
       [ 8, 10]])>
```

## Import the necessary libraries

```
[97]:  
import datetime, warnings, scipy  
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
plt.rcParams["patch.force_edgecolor"] = True  
plt.style.use('fivethirtyeight')  
mpl.rc('patch', edgecolor='dimgray', linewidth=1)  
from IPython.core.interactiveshell import InteractiveShell  
InteractiveShell.ast_node_interactivity = "last_expr"  
pd.options.display.max_columns = 50  
%matplotlib inline  
warnings.filterwarnings("ignore")
```

## Import the functions from Keras and sklearn libraries

```
[106]:  
import math  
from keras.models import Sequential  
from keras.layers import Dense  
from keras.layers import LSTM  
from sklearn.preprocessing import MinMaxScaler, StandardScaler  
from sklearn.metrics import mean_squared_error
```

# Notes by keras application and me

You will also need to install a backend framework – either JAX, TensorFlow, or PyTorch:

And

```
pip install --upgrade keras-cv
```

```
pip install --upgrade keras-nlp
```

```
pip install --upgrade keras
```

```
[4]: import os  
os.environ["KERAS_BACKEND"] = "jax"  
import keras
```

Note: The backend must be configured before importing Keras, and the backend cannot be changed after the package has been imported.

```
[5]: import os  
os.environ["TF_USE_LEGACY_KERAS"] = "1"  
import tensorflow
```

Note: These line would need to be before any import tensorflow

## ▶ Introduction

Keras 3 is a deep learning framework works with TensorFlow, JAX, and PyTorch interchangeably. This notebook will walk you through key Keras 3 workflows.

Keras, a popular deep learning library in Python, is divided into several key modules:

**Layers API:** Provides various types of neural network layers, like Dense, Conv2D, LSTM, etc., which can be used to build the architecture of a neural network.

**Models API:** Includes the Sequential model for building simple linear stacks of layers and the Functional API for creating more complex models.

**Losses API:** Offers various loss functions like `mean_squared_error`, `binary_crossentropy`, etc., which are used to measure the difference between the predicted and actual outputs.

**Optimizers API:** Contains different optimization algorithms like SGD, Adam, RMSprop, etc., which are used to adjust the weights of the neural network to minimize the loss function.

**Metrics API:** Provides different metrics like accuracy, precision, recall, etc., to evaluate the performance of a model.

**Callbacks API:** Enables the use of custom actions at different stages of training, such as early stopping, learning rate scheduling, etc.

**Datasets API:** Includes pre-built datasets like MNIST, CIFAR-10, IMDB, etc., for training and evaluating models.

**Utilities API:** Contains utility functions for various tasks like image preprocessing, text tokenization, and saving/loading models.

```
# Here are some common Keras syntax examples for different tasks:

1. Building a Sequential Model:
from keras.models import Sequential
from keras.layers import Dense

# Initialize the model
model = Sequential()

# Add layers
model.add(Dense(units=64, activation='relu', input_shape=(input_dim,)))
model.add(Dense(units=10, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

2. Building a Model using the Functional API:
from keras.models import Model
from keras.layers import Input, Dense

# Define the input layer
inputs = Input(shape=(input_dim,))

# Add layers
x = Dense(64, activation='relu')(inputs)
outputs = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=inputs, outputs=outputs)

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
3. Training a Model:  
# Train the model  
model.fit(x_train, y_train, epochs=10, batch_size=32, validation_split=0.2)  
4. Evaluating a Model:  
# Evaluate the model  
loss, accuracy = model.evaluate(x_test, y_test, batch_size=32)  
print(f'Test loss: {loss}')  
print(f'Test accuracy: {accuracy}')  
5. Making Predictions:  
# Make predictions  
predictions = model.predict(x_new)  
6. Using Callbacks:  
from keras.callbacks import EarlyStopping  
  
# Define early stopping  
early_stopping = EarlyStopping(monitor='val_loss', patience=3)  
  
# Train the model with callbacks  
model.fit(x_train, y_train, epochs=10, batch_size=32, validation_split=0.2, callbacks=[early_stopping])  
7. Saving and Loading a Model:  
# Save the model  
model.save('my_model.h5')  
  
# Load the model  
from keras.models import load_model  
model = load_model('my_model.h5')  
These examples cover the basics of defining, training, evaluating, and saving models in Keras.
```

```
[2]:
```

```
import numpy as np
import os

os.environ["KERAS_BACKEND"] = "jax"

# Note that Keras should only be imported after the backend
# has been configured. The backend cannot be changed once the
# package is imported.

import keras
```

# Note that Keras should only be imported after the backend  
# has been configured. The backend cannot be changed once the  
# package is imported.

A first example: A MNIST convnet Let's start with the Hello World of ML: training a convnet to classify MNIST digits.

```
[3]:
```

```
# Load the data and split it between train and test sets
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Scale images to the [0, 1] range
x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255
# Make sure images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
print(x_train.shape[0], "train samples")
print(x_test.shape[0], "test samples")
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

**11490434/11490434** **13s** 1us/step  
x\_train shape: (60000, 28, 28, 1)  
y\_train shape: (60000,)  
60000 train samples  
10000 test samples

Certainly! Here's a step-by-step explanation of what each part of the provided code does:

1. Load the Data

```
python
```

[Copy code](#)

```
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

**What It Does:** This line loads the MNIST dataset from Keras. MNIST is a collection of 28x28 grayscale images of handwritten digits (0-9).

**Outputs:**

`x_train`: The training images (60,000 images).

`y_train`: The labels corresponding to the training images (60,000 labels).

`x_test`: The test images (10,000 images).

`y_test`: The labels corresponding to the test images (10,000 labels).

2. Scale Images to the [0, 1] Range

```
python
```

[Copy code](#)

```
x_train = x_train.astype("float32") / 255
```

```
x_test = x_test.astype("float32") / 255
```

**What It Does:** This normalizes the pixel values of the images. Originally, pixel values range from 0 to 255. Dividing by 255 scales these values to the range [0, 1], which is more suitable for neural network training.

**Result:** Each pixel value in the images is now a float between 0 and 1.

### 3. Ensure Images Have Shape (28, 28, 1)

python

Copy code

```
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
```

What It Does: Adds an extra dimension to the image arrays to represent the number of color channels. Since MNIST images are grayscale, this dimension will be 1.

Result:

The shape of x\_train changes from (60,000, 28, 28) to (60,000, 28, 28, 1).

The shape of x\_test changes from (10,000, 28, 28) to (10,000, 28, 28, 1).

### 4. Print Shapes and Number of Samples

python

Copy code

```
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
print(x_train.shape[0], "train samples")
print(x_test.shape[0], "test samples")
```

What It Does: Prints the shapes of the image and label arrays and the number of samples in the training and test sets.

Outputs:

x\_train shape: Shows the shape of the training image array (should be (60000, 28, 28, 1)).

y\_train shape: Shows the shape of the training labels array (should be (60000,)).

x\_train.shape[0]: Prints the number of training samples (60,000).

x\_test.shape[0]: Prints the number of test samples (10,000).

```
[7]: # Model parameters
num_classes = 10
input_shape = (28, 28, 1)

model = keras.Sequential(
    [
        keras.layers.Input(shape=input_shape),
        keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        keras.layers.MaxPooling2D(pool_size=(2, 2)),
        keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
        keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
        keras.layers.GlobalAveragePooling2D(),
        keras.layers.Dropout(0.5),
        keras.layers.Dense(num_classes, activation="softmax"),
    ]
)
```

```
[8]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 64)	640
conv2d_1 (Conv2D)	(None, 24, 24, 64)	36,928
max_pooling2d (MaxPooling2D)	(None, 12, 12, 64)	0
conv2d_2 (Conv2D)	(None, 10, 10, 128)	73,856
conv2d_3 (Conv2D)	(None, 8, 8, 128)	147,584
global_average_pooling2d (GlobalAveragePooling2D)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 10)	1,290

Total params: 260,298 (1016.79 KB)

Trainable params: 260,298 (1016.79 KB)

Non-trainable params: 0 (0.00 B)

The code snippet you provided defines a Convolutional Neural Network (CNN) using Keras. Here's a detailed step-by-step explanation of each part:

1. Define Model Parameters

Python

Copy code

```
num_classes = 10
input_shape = (28, 28, 1)

num_classes: Number of output classes for classification. MNIST has 10 classes (digits 0-9).
input_shape: Shape of the input images, which are 28x28 pixels with 1 color channel (grayscale).
```

2. Define the Model

Python

Copy code

```
model = keras.Sequential(
```

```
[    keras.layers.Input(shape=input_shape),
    keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
    keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),
    keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
    keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
    keras.layers.GlobalAveragePooling2D(),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(num_classes, activation="softmax"),
]
```

## Explanation of Each Layer:

### Input Layer:

```
python  
Copy code
```

```
keras.layers.Input(shape=input_shape)
```

Purpose: Specifies the shape of the input data (28x28 pixels with 1 color channel). This is the first layer of the model.

### First Convolutional Layer:

```
python  
Copy code
```

```
keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu")
```

Purpose: Applies 64 convolutional filters of size 3x3 to the input image.

Activation Function: ReLU (Rectified Linear Unit) introduces non-linearity.

### Second Convolutional Layer:

```
python  
Copy code
```

```
keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu")
```

Purpose: Applies another set of 64 convolutional filters of size 3x3.

Activation Function: ReLU.

### Max Pooling Layer:

```
python  
Copy code
```

```
keras.layers.MaxPooling2D(pool_size=(2, 2))
```

Purpose: Reduces the spatial dimensions of the feature maps by taking the maximum value in each 2x2 block.  
Third Convolutional Layer:

```
python  
Copy code
```

```
keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu")
```

Purpose: Applies 128 convolutional filters of size 3x3.

Activation Function: ReLU.

Fourth Convolutional Layer:

```
python  
Copy code
```

```
keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu")
```

Purpose: Applies another set of 128 convolutional filters of size 3x3.

Activation Function: ReLU.

Global Average Pooling Layer:

```
python  
Copy code
```

```
keras.layers.GlobalAveragePooling2D()
```

Purpose: Reduces each feature map to a single value by averaging over all spatial dimensions. This flattens the output from the previous layer while preserving the number of feature maps.

Dropout Layer:

```
python  
Copy code
```

```
keras.layers.Dropout(0.5)
```

Purpose: Applies dropout regularization with a rate of 0.5, randomly setting 50% of the input units to 0 during training to prevent overfitting.  
Dense Output Layer:

```
python  
Copy code
```

```
keras.layers.Dense(num_classes, activation="softmax")
```

Purpose: A fully connected layer with num\_classes (10) units. Each unit represents a class probability.

Activation Function: Softmax converts the output into a probability distribution over the 10 classes.

```
[9]: # We use the compile() method to specify the optimizer, loss function,  
# and the metrics to monitor. Note that with the JAX and TensorFlow  
# backends, XLA compilation is turned on by default.
```

```
model.compile(  
    loss=keras.losses.SparseCategoricalCrossentropy(),  
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),  
    metrics=[  
        keras.metrics.SparseCategoricalAccuracy(name="acc"),  
    ],  
)
```

```
[*]: # Let's train and evaluate the model.  
# We'll set aside a validation split of 15% of the data during training  
# to monitor generalization on unseen data.
```

```
batch_size = 128  
epochs = 20  
  
callbacks = [  
    keras.callbacks.ModelCheckpoint(filepath="model_at_epoch_{epoch}.keras"),  
    keras.callbacks.EarlyStopping(monitor="val_loss", patience=2),  
]  
  
model.fit(  
    x_train,  
    y_train,  
    batch_size=batch_size,  
    epochs=epochs,  
    validation_split=0.15,  
    callbacks=callbacks,  
)  
score = model.evaluate(x_test, y_test, verbose=0)
```

The provided code snippet outlines the process to train and evaluate the CNN model. Here's a detailed step-by-step explanation of each part:

1. Set Training Parameters

python

[Copy code](#)

```
batch_size = 128
```

```
epochs = 20
```

**batch\_size:** The number of samples processed before the model is updated. Here, 128 samples are used per batch.

**epochs:** The number of times the entire training dataset will pass through the model. Here, the model will train for 20 epochs.

2. Define Callbacks

python

[Copy code](#)

```
callbacks = [
```

```
    keras.callbacks.ModelCheckpoint(filepath="model_at_epoch_{epoch}.keras"),
```

```
    keras.callbacks.EarlyStopping(monitor="val_loss", patience=2),
```

```
]
```

**ModelCheckpoint:** Saves the model's weights at the end of each epoch. The filepath specifies where to save the model and includes {epoch} to save different files for each epoch.

**EarlyStopping:** Stops training early if the validation loss does not improve for a specified number of epochs (patience=2). This helps prevent overfitting and reduces unnecessary computation.

### 3. Train the Model

```
python  
Copy code  
model.fit(  
    x_train,  
    y_train,  
    batch_size=batch_size,  
    epochs=epochs,  
    validation_split=0.15,  
    callbacks=callbacks,  
)
```

x\_train and y\_train: The training data and labels.

batch\_size: The number of samples processed before updating the model weights.

epochs: The total number of epochs to train the model.

validation\_split: A fraction of the training data (15%) used for validation during training. This helps monitor the model's performance on unseen data while training.

callbacks: List of callbacks to apply during training. Here, it includes saving the model and early stopping.

### 4. Evaluate the Model

```
python  
Copy code
```

```
score = model.evaluate(x_test, y_test, verbose=0)
```

x\_test and y\_test: The test data and labels.

verbose=0: Suppresses detailed output during evaluation.

```
[*]: # During training, we were saving a model at the end of each epoch.  
      # You can also save the model in its latest state like this:  
      model.save("final_model.keras")  
  
      # And reload it like this:  
      model = keras.saving.load_model("final_model.keras")  
  
      # Next, you can query predictions of class probabilities with predict():  
      predictions = model.predict(x_test)
```